

Assignment

Logical Volume Manager (LVM) is a disk management technology in Linux that allows you to manage disk drives and storage volumes with greater flexibility compared to traditional partitioning methods.

Key Concepts:

1) Physical Volume (PV): Represents a physical disk or a partition that LVM can use.

2) Volume Group (VG): A pool of storage created by grouping multiple physical volumes.

3) Logical Volume (LV): The equivalent of a partition created from the space within a volume group, which can be resized as needed.

To create Volumes for the drive, let's say three:

`pvcreeate /dev/sdb /dev/sdc /dev/sdd`

Before this command, we need to add three disks. You are initializing three physical disks (/dev/sdb, /dev/sdc, and /dev/sdd) as Physical Volumes (PVs) in LVM.

`pvddisplay`

The `pvddisplay` command in LVM is used to display detailed information about physical volumes. After creating physical volumes with `pvcreeate`, you can use `pvddisplay` to get detailed information like the size, volume group association, status, and more.

```
[root@localhost ~]# pvcreeate /dev/sdb /dev/sdc /dev/sdd
Physical volume "/dev/sdb" successfully created.
Physical volume "/dev/sdc" successfully created.
Physical volume "/dev/sdd" successfully created.
[root@localhost ~]# pvddisplay
--- Physical volume ---
PV Name           /dev/sda2
VG Name           cs
PV Size           <19.00 GiB / not usable 3.00 MiB
Allocatable       yes (but full)
PE Size           4.00 MiB
Total PE          4863
Free PE           0
Allocated PE       4863
PV UUID           H150Ea-PTBv-8Pb0-3vHK-P1nY-g5r0-Q1XMox

"/dev/sdb" is a new physical volume of "20.00 GiB"
--- NEW Physical volume ---
PV Name           /dev/sdb
VG Name
```

vgcreate techmpst /dev/sdb /dev/sdc /dev/sdd

It creates a Volume Group (VG) named techmpst by combining the three physical volumes /dev/sdb, /dev/sdc, and /dev/sdd. This volume group pools the storage from these disks, allowing you to create flexible logical volumes (LVs) on top of it.

vgdisplay

The vgdisplay command in LVM provides detailed information about volume groups, including their name, size, free space, and physical extent details. It displays whether the volume group is allocatable, the number of physical extents (PEs), and the unique identifier (UUID) of the volume group. This helps in monitoring and managing storage pools created from physical volumes.

```
[root@localhost ~]# vgcreate techmpst /dev/sdb/ /dev/sdc /dev/sdd
No device found for /dev/sdb/.
Command requires all devices to be found.
[root@localhost ~]# vgcreate techmpst /dev/sdb /dev/sdc /dev/sdd
Volume group "techmpst" successfully created
[root@localhost ~]# vgdisplay
--- Volume group ---
VG Name                techmpst
System ID
Format                 lvm2
Metadata Areas         3
Metadata Sequence No   1
VG Access               read/write
VG Status               resizable
MAX LV                 0
Cur LV                 0
Open LV                 0
Max PV                 0
Cur PV                 3
Act PV                 3
VG Size                 <59.99 GiB
PE Size                 4.00 MiB
Total PE                15357
Alloc PE / Size         0 / 0
Free PE / Size          15357 / <59.99 GiB
VG UUID                 BHWLW6-kWK7-UcHg-YbVr-89Jo-cJyu-drqt3H

--- Volume group ---
VG Name                cs
```

lvcreate -L 8000 techmpst -n sales /dev/sdb /dev/sdc /dev/sdd

It is used to create a Logical Volume (LV) named sales with a size of 8000MB (8GB) from the Volume Group (VG) named techmpst. The -L 8000 option specifies the desired size for the logical volume, while the -n sales option sets its name. Importantly, there is no need to specify the individual physical volumes (such as /dev/sdb, /dev/sdc, or /dev/sdd) again, as they are already included in the techmpst volume group. After executing this command, the

logical volume can be formatted with a filesystem and mounted for use in your system.

In the same manner,

```
lvcreate -L 8000 techmpst -n mkt /dev/sdb /dev/sdc /dev/sdd
```

```
lvcreate -L 4000 techmpst -n IT /dev/sdb /dev/sdc /dev/sdd
```

```
lvcreate -L 4000 techmpst -n researchdept /dev/sdb /dev/sdc /dev/sdd
```

```
[root@localhost ~]# lvcreate -L 8000 techmpst -n sales /dev/sdb /dev/sdc /dev/sdd
Logical volume "sales" created.
[root@localhost ~]# lvcreate -L 8000 techmpst -n mkt /dev/sdb /dev/sdc /dev/sdd
Logical volume "mkt" created.
[root@localhost ~]# lvcreate -L 4000 techmpst -n IT /dev/sdb /dev/sdc /dev/sdd
Logical volume "IT" created.
```

```
[root@localhost ~]# lvcreate -L 4000 techmpst -n researchdept /dev/sdb /dev/sdc /dev/sdd
Logical volume "researchdept" created.
```

lvdisplay

It is used to show detailed information about logical volumes (LVs) in LVM (Logical Volume Manager). When executed, it provides a summary of each logical volume, including its name, the volume group it belongs to, the device path, size, number of allocated and current logical extents, type, access permissions, status, creation time, and unique identifier (UUID).

```
[root@localhost ~]# lvdisplay
--- Logical volume ---
LV Path                /dev/techmpst/sales
LV Name                 sales
VG Name                techmpst
LV UUID                e0Cr6l-HXNu-5Z1L-zC25-A50t-CXBc-YVRla4
LV Write Access         read/write
LV Creation host, time localhost.localdomain, 2024-09-24 17:47:25 +0530
LV Status               available
# open                  0
LV Size                 7.81 GiB
Current LE              2000
Segments                1
Allocation              inherit
Read ahead sectors     auto
 - currently set to    256
Block device            253:3

--- Logical volume ---
LV Path                /dev/techmpst/mkt
LV Name                mkt
VG Name                techmpst
LV UUID                76hGzH-7vqI-T3bJ-u8Ee-5ePV-bse3-aLXAex
LV Write Access         read/write
LV Creation host, time localhost.localdomain, 2024-09-24 17:47:38 +0530
LV Status               available
# open                  0
LV Size                 7.81 GiB
Current LE              2000
Segments                1
Allocation              inherit
Read ahead sectors     auto
```

mkdir /sales /mkt /IT /researchdept

It creates four directories in the root filesystem: /sales, /mkt, /IT, and /researchdept. Each directory serves as a separate folder for organizing files and data related to their respective departments or purposes.

Breakdown of the Command:

mkdir: Stands for "make directory," which is the command used to create new directories.

/sales: Creates a directory named sales.

/mkt: Creates a directory named mkt (possibly for marketing).

/IT: Creates a directory named IT (for information technology).

/researchdept: Creates a directory named researchdept (for the research department).

mount -a /techmpst /sales /sales

Breakdown of the Commands:

mount -a:

This command mounts all filesystems listed in the /etc/fstab file that are not currently mounted. It does not require any additional parameters.

mount /dev/techmpst /sales:

This command mounts the logical volume located at /dev/techmpst to the directory /sales. Make sure that the /sales directory exists before you attempt to mount.

The same manner,

mount -a /techmpst /mkt /mkt

mount -a /techmpst /IT /IT

mount -a /techmpst /researchdept / researchdept

mount

This command is used to attach filesystems to a specified mount point in the directory tree. When you run the mount command without any arguments, it displays a list of all currently mounted filesystems along with their mount points and other details.

RAID:

RAID (Redundant array of independent disks) is a storage technology that combines multiple disk drives into one logical unit. It can help improve performance and data redundancy. CentOS, a Linux distribution, supports RAID configurations during installation and offers a few ways to use RAID.

Categories of RAID:

RAID 0:

RAID 0, or “striping,” is a RAID configuration involving the division of data across the disks where there is no redundancy of the disks. It increases performance since it provides both read and write access, while it does not support data backup. Data is only safe when it’s stored in two drives at once; if one fails, then the data is gone. RAID 0 is used for applications that demand very high speed but do not concern themselves with data protection.

RAID 1:

RAID 1 is defined as the ‘mirroring’ in which the same set of data is written onto two or more drives. It will help to have duplication, and if one of the drives fails, the data is still available in the other drives. Although it provides extremely high data reliability and data protection features, this concept does not enhance performance and demands double space as compared to the data.

RAID 2:

It consist of Bit-level Striping. RAID 2 records Error Correction Code (ECC) using hamming code parity. In this level each data bit in a word is recorded on a separate disk and ECC codes of the data words are stored on a different set of disk.

Advantages–

1. In case of Error Correction it uses hamming code.
2. It Uses one designated drive to store parity.

Disadvantage –

1. It has complex structure and high cost due to extra drive.
2. It require extra drive for error detection.

RAID 3:

RAID 3 consists of Byte-level Striping. It stripes the data onto multiple disk. The parity bit generated for each disk section and stored on a different dedicated disk. This level overcome the single disk failure.

Advantages –

1. Data can be transferred in bulk.
2. Data can be accessed in parallel.

Disadvantages –

1. It require an additional drive for parity.
2. In case of small size files it performs slowly

RAID 4:

Block-level striping and a separate parity disk are used in RAID 4. This configuration stores the parity information for each data block on a different disk, and divides the data into blocks that are striped across multiple disks.

RAID-5 (Block-Level Stripping with Distributed Parity):

This is a slight modification of the RAID-4 system where the only difference is that the parity rotates among the drives.

RAID-6 (Block-Level Striping with two Parity Bits):

Raid-6 helps when there is more than one disk failure. A pair of independent parities are generated and stored on multiple disks at this level. Ideally, you need four disk drives for this level.

There are also hybrid RAIDs, which make use of more than one RAID level nested one after the other, to fulfill specific requirements

Implementation:

Initially, remove the logical volume and volume group named techmpst:

lvremove techmpst

vgremove techmpst

Then, remove the physical volume:

pvremove /dev/sdb /dev/sdc /dev/sdd

mdadm -C /dev/md0 -n3 /dev/sdb /dev/sdc /dev/sdd -l5

It is used to create a RAID 5 array. Let's break it down:

mdadm: This is the command-line tool used for managing RAID arrays.

-C: This option specifies that we want to create a new array.

/dev/md0: This is the name of the RAID array we're creating. It will be accessible as a block device.

-n3: This option specifies the number of disks (or devices) that will be part of the RAID array. In this case, it's 3.

/dev/sdb /dev/sdc /dev/sdd: These are the paths to the physical disks that will be used to create the RAID array.

-l5: This option specifies the RAID level. In this case, it's RAID 5.

```

[root@localhost ~]# lvremove techmpst
Do you really want to remove active logical volume techmpst/sales? [y/n]: y
Logical volume "sales" successfully removed.
Do you really want to remove active logical volume techmpst/mkt? [y/n]: y
Logical volume "mkt" successfully removed.
Do you really want to remove active logical volume techmpst/IT? [y/n]: y
Logical volume "IT" successfully removed.
Do you really want to remove active logical volume techmpst/researchdept? [y/n]: y
Logical volume "researchdept" successfully removed.
[root@localhost ~]# vgremove techmpst
Volume group "techmpst" successfully removed
[root@localhost ~]# pvremove /dev/sdb /dev/sdc /dev/sdd
Labels on physical volume "/dev/sdb" successfully wiped.
Labels on physical volume "/dev/sdc" successfully wiped.
Labels on physical volume "/dev/sdd" successfully wiped.
[root@localhost ~]# mdadm -C /dev/md0 -n3 /dev/sdb /dev/sdc /dev/sdd -l5
mdadm: Defaulting to version 1.2 metadata
mdadm: array /dev/md0 started.

```

mdadm -D /dev/md0

It is used to display detailed information about a RAID array named /dev/md0

```

[root@localhost ~]# mdadm -D /dev/md0
/dev/md0:
    Version : 1.2
    Creation Time : Tue Sep 24 21:41:35 2024
    Raid Level : raid5
    Array Size : 41908224 (39.97 GiB 42.91 GB)
    Used Dev Size : 20954112 (19.98 GiB 21.46 GB)
    Raid Devices : 3
    Total Devices : 3
    Persistence : Superblock is persistent

    Update Time : Tue Sep 24 21:42:00 2024
    State : clean, degraded, recovering
    Active Devices : 2
    Working Devices : 3
    Failed Devices : 0
    Spare Devices : 1


    Layout : left-symmetric
    Chunk Size : 512K
Consistency Policy : resync

    Rebuild Status : 21% complete

    Name : localhost.localdomain:0 (local to host localhost.localdomain)
    UUID : c8b2c92d:face6182:0d6b4666:df0d9ead
    Events : 4

   Number   Major   Minor   RaidDevice State
    -----   -----   -----   -
    0         8       16         0      active sync   /dev/sdb
    1         8       32         1      active sync   /dev/sdc
    3         8       48         2      spare rebuilding /dev/sdd

```

mkdir /raid

It creates a new directory named "raid" in the root directory of your filesystem.

mkfs.ext4 /dev/md0

It is used to format a RAID array named /dev/md0 with the EXT4 filesystem.

```
[root@localhost ~]# mkfs.ext4 /dev/md0
mke2fs 1.46.5 (30-Dec-2021)
Creating filesystem with 10477056 4k blocks and 2621440 inodes
Filesystem UUID: 39120905-4ead-484a-adc1-050d38def808
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
    4096000, 7962624

Allocating group tables: done
Writing inode tables: done
Creating journal (65536 blocks): done
Writing superblocks and filesystem accounting information: done

[root@localhost ~]# mkfs.xfs -f /dev/md0
log stripe unit (524288 bytes) is too large (maximum is 256KiB)
log stripe unit adjusted to 32KiB
meta-data=/dev/md0            isize=512    agcount=16, agsize=654720 blks
                               sectsz=512    attr=2, projid32bit=1
                               crc=1          finobt=1, sparse=1, rmapbt=0
                               reflink=1      bigtime=1 inobtcount=1 nrext64=0
data      =                   bsize=4096   blocks=10475520, imaxpct=25
                               sunit=128     swidth=256 blks
naming    =version 2          bsize=4096   ascii-ci=0, ftype=1
log       =internal log      bsize=4096   blocks=16384, version=2
                               sectsz=512    sunit=8 blks, lazy-count=1
realtime  =none              extsz=4096   blocks=0, rtextents=0
```

mount /dev/md0 /raid

It mounts the RAID array /dev/md0 to the directory /raid, making its contents accessible.

cd /raid

It is used to change the current directory to /raid.

touch t1 t2 t3

It is used to create three empty files named t1, t2, and t3 in the current directory i.e raid.

To remove HDD if disk is faulty:

mdadm -f /dev/md0 /dev/sdd

It marks the disk /dev/sdd as failed in the RAID array /dev/md0, triggering a rebuild process.

mdadm -r /dev/md0 /dev/sdd

It starts a rebuild process for the failed disk /dev/sdd in the RAID array /dev/md0.

```
[root@localhost raid]# mdadm -f /dev/md0 /dev/sdd
mdadm: set /dev/sdd faulty in /dev/md0
[root@localhost raid]# mdadm -r /dev/md0 /dev/sdd
mdadm: hot removed /dev/sdd from /dev/md0
```

mdadm -a /dev/md0 /dev/sdd

It is used to add a disk to a RAID array.

```
[root@localhost raid]# mdadm -a /dev/md0 /dev/sdd
mdadm: added /dev/sdd
```

mdadm -S /dev/md0 /dev/sdd

It is used to stop the rebuild process for a specific disk in a RAID array.